

Eventless Payload-Adaptive NMPC for Quadrotors

Anonymous Author(s)

Affiliation

[TODO: authors, affiliation, email]

Abstract—A release command does not guarantee physical payload detachment, so switching a controller’s model from an attach/drop signal can be unsafe. We present an eventless payload-adaptive control framework in which a moving horizon estimator (MHE) uses only motor speeds and odometry to publish an atomic, health-gated frame of composite mass, first mass moment $s = m_{prxy}$, and inertia to an outer-loop nonlinear model predictive controller (NMPC). The first-moment representation remains well defined as the payload vanishes, while bounded filtering and continuous body-rate scheduling prevent discontinuous model updates. A thrust-residual self-trigger deweights stale pre-change stages in the MHE window and reduces mass-estimate settling time by up to $4\times$. In PX4 software-in-the-loop experiments with physically attached eccentric payloads, the eventless interface completes release confirmation in 22/22 valid paired flights, with no observed post-release instability and a median confirmation delay of 2.70 s. A safety ablation shows that payload quantity and thrust residual are individually unreliable during aggressive flight: their fast release branch causes premature model clearing in 3/9 flights. Requiring fresh moment-collapse evidence eliminates premature declarations in a 12-flight smoke test, yielding 11 confirmations and one conservative unresolved outcome. Against an L1-adaptive NMPC baseline, the proposed method provides comparable transient tracking, recovers 15% to 35% faster, and retains interpretable physical parameters. These results demonstrate the eventless control interface and its fail-safe semantics, while not claiming distribution-wide robustness of the release detector.

I. INTRODUCTION

Autonomous aerial delivery and in-flight aerial manipulation subject a multirotor to abrupt changes in its own dynamics. Grasping or releasing a parcel changes three coupled quantities at once: the total *mass* m (which rescales the thrust-to-acceleration map and the hover throttle), the *center of mass* (a thrust that no longer passes through the composite CoM produces a constant trim torque $\tau \approx \Delta r \times f$, significant even for small offsets because $f \approx mg$ is large), and the *rotational inertia* (which slows the attitude response and effectively lowers the inner-loop bandwidth). The first channel is an outer-loop (position/thrust) problem; the latter two are inner-loop (attitude/rate) problems. A controller that models the payload only as a scalar mass will mistrack or, for sufficiently heavy eccentric loads, destabilize.

Three families of methods address load changes. *Explicit parameter estimation* (EKF/MHE of mass, CoM, inertia) is interpretable, integrates natively with a model-based controller, and can update thrust-margin constraints, but classical estimators converge slowly across an abrupt change because the estimation window or filter averages across the discontinuity. *Lumped-disturbance adaptation* (L1-adaptive NMPC, INDI)

compensates the aggregate effect without estimating physical parameters, reacting fast but yielding no interpretable model. *Learning-based* methods (e.g. neural MHE that regresses estimator weights) can resolve the estimator’s bias–variance trade-off but typically require rich sensing or offline ground-truth data.

The central question of this work is whether the complete grasp–carry–release lifecycle can adapt its payload model without trusting a commanded or externally supplied payload event. Our answer rests on two observations. First, an abrupt payload change leaves a detectable common-mode signature in reconstructed collective thrust, so the estimator can discover *when* the dynamics changed without consuming a gripper command, attachment state, or external event notification. Second, **motor speeds are a free physical sensor**: both collective thrust and body torques can be reconstructed from rotor speeds alone, yielding physical residuals that are decoupled from the estimator’s own state. Combining the two yields an estimator that is both fast (it deweights stale pre-event measurements the moment the thrust residual crosses a threshold) and *sensing-minimal* (it needs no external event signal and no measurement of the attachment geometry).

Contributions.

- 1) An **eventless payload-adaptive NMPC architecture** whose model never consumes an attach/drop notification or treats its own release command as evidence. An atomic, health-gated $m/s/J$ frame drives bounded model updates and continuous body-rate scheduling (§III–IV).
- 2) A **self-triggered MHE** that detects abrupt payload changes from reconstructed thrust and deweights stale pre-event window stages, reducing mass-estimate settling time by up to $\sim 4\times$ without an external event signal. The first-moment state $s = m_{prxy}$ represents eccentricity without measured attachment geometry and remains well-defined at release (§IV; §VI-A).
- 3) **Closed-loop lifecycle and safety validation**: 22/22 valid paired flights complete eventless release without observed post-release instability. A targeted ablation rejects an unsafe quantity-plus-residual branch and shows that fresh moment-collapse evidence with an UNRESOLVED fail-safe state is necessary (§VI-C).

The runtime estimator consumes only motor speeds and odometry; the sole geometric quantity not identified online is a weak vertical-arm prior r_z , which is structurally unobservable from the thrust-moment channel.

II. RELATED WORK

Control under changing payload. Adaptive and robust designs—L1-adaptive NMPC [1], INDI [2]—treat the load as a lumped disturbance and are the standard baselines for payload-carrying flight; L1-NMPC in particular reports large tracking-error reductions under unknown load on aggressive trajectories. Slung-load control addresses the swinging-payload case explicitly. These methods do not produce a physical parameter and cannot, e.g., update a thrust-margin constraint from the estimate.

MHE and online parameter estimation. Filter- and optimization-based estimators recover mass, inertia, and geometric parameters [3]–[5]. Standard MHE performance is highly sensitive to the process/measurement weight matrices; neural MHE (NeuroMHE) [6] learns to output these weights online, resolving the bias–variance trade-off but estimating a lumped disturbance force rather than explicit named parameters.

Event-triggered and change detection. Residual-based fault/change detection (CUSUM and related) underlies our self-trigger. Our positioning: we do not learn dynamics and add no sensing; every deployed signal derives from motor speeds and odometry. We separately test, as a supporting ablation, whether numerical tuning improves the handcrafted transient-period schedule.

III. SYSTEM OVERVIEW AND PROBLEM SETUP

Architecture (Fig. 1). Three layers: Gazebo physics (with a magnetic DetachableJoint gripper producing genuine inertia/CoM changes), the PX4 inner-loop rate controller, and a ROS 2 outer loop (NMPC + MHE). The outer loop publishes a collective thrust and a body-rate setpoint; PX4 tracks the body rate.

Dynamics and parameterization. The NMPC model carries the runtime parameter vector $p = [x_{\text{ref}}, m, \Delta J, c_{xy}]$ (extended with a lumped disturbance only for the L1 baseline). Thrust acting at the geometric center about the composite CoM produces $\tau_{\text{com}} = [-c_y T, +c_x T, 0]$. The deployed MHE augments both composite mass and the horizontal first mass moment,

$$x_{\text{MHE}} = [x; m; s_x; s_y] \in \mathbb{R}^{16}, \quad s = m p r_{xy}, \quad (1)$$

with $\dot{m} = 0$ and $\dot{s} = 0$ within each window ($N = 20$, $\Delta t = 0.1$ s). The published geometry is reconstructed as

$$c_{xy} = \frac{s}{m}, \quad \mu = \frac{m_B \max(m - m_B, 0)}{m}, \quad (2)$$

$$\Delta J_{xx} = \Delta J_{yy} = \mu r_z^2.$$

where r_z is a weak rack prior. The small $O(r_{xy}^2)$ diagonal terms and yaw increment are omitted in the deployed interface. The hover-input baseline is recomputed from the live model and includes the eccentric trim torque, avoiding a stale operating-point bias.

Model-independent physical inputs. We reconstruct the collective thrust $T_{\text{phys}} = k \sum_i \omega_i^2$ and the body torques from rotor speeds and arm geometry. Feeding the NMPC’s

commanded input back into the estimator instead creates a self-consistent blind spot: at hover $m \approx T/g$ is self-consistent for any m , so the mass is unobservable from the commanded thrust. The motor-speed reconstruction breaks this because it is decoupled from the controller/estimator state.

Thrust-to-throttle inversion. The offboard collective-thrust command is inverted to a normalized rotor command through the motor model $T = k\omega^2$; because k appears on both the forward and inverse paths, steady-state tracking is independent of the absolute thrust calibration.

Continuous inner-loop scheduling. Model consistency in the NMPC cannot by itself restore the PX4 rate loop’s lost bandwidth. The deployed path therefore maps the live inertia increment to a body-rate-command scale, low-pass filters it, applies target hysteresis and asymmetric rise/recovery slew limits, and preserves 5% command headroom before the final saturation. It does *not* rewrite PX4 parameters. An attach-time inertia-only bootstrap uses the rack’s 0.30 kg payload envelope for immediate safety, then fades out after persistent coherent MHE evidence; neither payload mass nor s is injected by this bootstrap, and the drop path remains estimate-only.

IV. SELF-TRIGGERED ESTIMATION AND EVENTLESS ADAPTATION

A. Deweighting mechanism

Within the sliding window, upon a payload event we deweight the R/Q blocks of the pre-event stages (per-stage `cost_set`), rather than softening the mass anchor Q_0 (which is already soft and is not the bottleneck; the true bottleneck is the $\dot{m} = 0$ rigidity that lets stale cross-discontinuity measurements pull the mass). Guardrails: the terminal measurement never enters the cost, and the first solve after an event keeps nominal weights.

B. No-signal self-trigger

The detector arms only after the baseline has settled (a warmup gate of consecutive quiet thrust-residual frames), then fires when $|T_{\text{phys}} - \text{baseline}|$ exceeds a threshold for a persistence count of consecutive frames. The warmup gate rejects the takeoff transient; the persistence count rejects gust pulses (broadband/transient) while catching a sustained mass step. This same criterion is the false-trigger criterion characterized under wind in §VI. This detector schedules the MHE window weights. Release confirmation uses the multi-source logic below; the two roles must not be conflated, because a threshold suitable for a 0.30 kg schedule event can suppress the residual vote of a 0.15 kg release.

C. Continuous payload interface and autonomous lifecycle

The MHE publishes one atomic frame

$$\mathcal{P} = \{m, s_x, s_y, c_x, c_y, \Delta J_{xx}, \Delta J_{yy}, \Delta J_{zz}, J_{xx}, J_{yy}, J_{zz}, \gamma, h, a\}, \quad (3)$$

where γ is no-payload confidence, h is a health bit and a is solution age. Health requires a successful latest solve and fresh odometry, motor-speed, and known-input streams. The




Fig. 1. System architecture and information flow. The full pipeline consumes only motor speeds and odometry; one atomic, health-gated $m/s/J$ frame drives both the NMPC model and continuous rate-command scheduling. No attach/drop notification enters this path.




Fig. 2. Sliding-window deweighting. On a payload event (thrust residual crossing αgm), the pre-event window stages are deweighted so stale cross-discontinuity measurements no longer pull the mass estimate.

NMPC additionally checks local receive age; an unhealthy or stale frame freezes the last accepted model and breaks the confidence persistence timer. Valid frames pass through a low-pass filter and separate slew limits for m , s , and ΔJ before every NMPC solve.

Payload presence is an estimator-owned state, deliberately distinct from the legacy flag that means “an external attach notification was received.” A hysteretic $\text{EMPTY} \leftrightarrow \text{LOADED}$ machine uses sustained payload-mass evidence, with a directional residual as a fast path, and a separate latch records that the present flight reliably entered LOADED . This latch prevents an empty-flight confidence value from consuming the later drop confirmation and prevents a transient mass dip during a figure-eight from disarming the true release path.

No-payload confidence is the product of smooth empty scores for payload quantity, first moment, and inertia. A nonzero moment therefore vetoes a false empty decision even if the mass estimate touches its lower bound. On accepted release evidence, the estimator switches only the *published target* of s to zero; the optimization state is not overwritten.

The output filter consequently produces a continuous decay rather than a geometric step. The controller marks a planned drop complete, resets its L1 state, and finishes recovery only after healthy confidence remains above 0.90 for 1.0 s; issuing the gripper command itself does none of these things.

D. Dimensionless trigger schedule

The deployed hand-designed schedule M0 applies a fixed pre-event deweighting rhythm and expresses the confirmation threshold as the dimensionless ratio α , with physical threshold αgm . This preserves the same relative margin across payload masses, whereas a fixed threshold in newtons does not. A supporting simulation-in-the-loop ablation tested whether optimizing the schedule magnitude improves M0; Sec. VI-B summarizes the negative result. M0 is retained in all reported system-level experiments.

E. Why the first mass moment is the essential state

Estimating CoM offset alone is ill-conditioned at release: physically $s = m_P r_{xy} \rightarrow 0$, whereas r_{xy} itself is undefined as $m_P \rightarrow 0$. In an earlier mass-only formulation, the optimizer could suppress a phantom eccentricity only by driving total mass below the bare-airframe value; 14 of 55 post-release samples then pinned at the lower bound. Augmenting s supplies the rotational channel directly and permits $c_{xy} = s/m$ without knowing the grasp-dependent horizontal offset. A frozen-window inertia coefficient prevents the optimizer from using instantaneous mass as a free inertia knob, while the weak r_z prior retains the dominant parallel-axis term. This structure is the minimum one that simultaneously represents payload existence, trim moment, and roll/pitch inertia without an external attachment geometry measurement.

The release detector studies three physically distinct sources: (i) collapse of $\rho = \|s\|/\|s\|_{\text{peak}}$, with an absolute $\|s\|$ bound used only as a sanity check; (ii) payload quantity, in which mass and algebraically derived inertia count as *one* source; and (iii) a directional thrust-residual edge with finite validity. A single source never releases. The slow path requires quantity-empty and moment-collapse evidence for five MHE

frames. The release-only residual is deliberately separate from the 4.4 N schedule-confirmation threshold and uses a signed two-window difference, $\Delta T = \bar{T}_{\text{post}} - \bar{T}_{\text{pre}}$, with 0.3 s half-windows. We evaluated two residual branches: fastA paired a normal 0.9 N residual with quantity-empty evidence, while fastB paired a longer-persistent residual with moment collapse. A recorded figure-eight transient reached -1.69 N, larger in magnitude than the -1.47 N step of the minimum supported payload, so amplitude alone cannot separate the cases. Online and offline evaluation call the same pure decision function. The randomized smoke of Sec. VI-C rejects fastA; the safety-accepted logic therefore requires moment-collapse evidence (slow or fastB) and treats an unresolved timeout as unknown, never as proof of an empty vehicle. Moment evidence must also originate from a fresh, successful MHE solve: the re-anchoring routine deliberately initializes s to zero after five consecutive solver failures, so that initialization is invalid evidence rather than a physical collapse.

V. SIMULATION SETUP

PX4 SITL with Gazebo (Harmonic), an x500 quadrotor, and a magnetic DetachableJoint gripper that produces genuine inertia and CoM changes.

Two payload-change mechanisms, by design. The physical DetachableJoint gripper attaches a separate rigid body and therefore couples mass, inertia and CoM together; it carries every system-level claim (Tables V, VI, and III). Separately, a *wrench-emulated* step—a virtual force injection that changes the effective translational mass instantaneously, with no inertia or CoM change—is used only where a controlled ablation demands it: isolating the event trigger’s estimation-layer benefit requires a mass discontinuity that is clean and repeatable, free of the grasp-depth scatter and inertia mismatch that a physical attach necessarily introduces (Table II and Fig. 3 only). The supporting weight-schedule ablation runs on the physical gripper, as do all baseline and wind results. The two mechanisms are complementary, not interchangeable, and we keep their scopes disjoint: any conclusion that needs genuine rigid-body coupling is drawn only from the physical experiments (see the red-line note in §VI-A). The wrench mechanism’s inability to change rotational inertia is a documented simulator limitation, not a modeling choice; we state it and its consequences in §VII.

Task state machine: approach (position control) → NMPC hand-off and warmup hover → controlled descent → active grasp enable → lift → figure-eight → phase-triggered release. We state the NMPC scope explicitly: it governs the grasp hover and the post-event adaptation phase; the approach controller is not part of the claim. Headless batch infrastructure with a GCS-heartbeat stand-in; the physical-effectiveness timestamp $t = 0$ is defined by T_{phys} crossing the baseline, and in-band is defined as ± 0.08 kg.

Frozen release configuration. In the evaluated implementation, `c_xy_est_enable` gates the update routine that contains both the legacy mass-domain test and the unified call to `release_decision()`. The node and generic launcher

TABLE I
MEASURED SOLVE TIMES AND LOOP RATES (SINGLE DESKTOP CORE, SHARED WITH GAZEBO/PX4 SITL). NMPC STATISTICS POOL $n = 1009$ SAMPLES ACROSS THE DEPLOYED 20 Hz BATCHES. MHE STATISTICS ARE FROM AN INSTRUMENTED HISTORICAL MASS-ONLY RUN ($n = 63$ REPORTING WINDOWS OF 20 SOLVES EACH), NOT FROM THE FINAL FIRST-MOMENT ESTIMATOR.

	NMPC	MHE	setpoint pub.
horizon / window	$N = 20, \Delta t = 0.05$ s (1.0 s preview)	$N = 20, \Delta t = 0.1$ s (2.0 s window)	—
state / control dim.	13 / 4	14 / — ^b	—
rate	20 Hz	10 Hz	50 Hz
period budget	50 ms	100 ms	20 ms
solve, median	1.3 ms	1.4 ms	—
solve, mean	1.35 ms	1.48 ms	—
solve, p95	2.1 ms	—	—
solve, p99	2.2 ms	—	—
solve, max	15.0 ms ^a	5.1 ms	—
median load	2.6 %	1.4 %	—
p99 load	4.4 %	—	—

^aSingle worst-case sample across all batches (first solve after solver warm-start); the p99 of 2.2 ms is the operative figure. ^bHistorical mass-only MHE; the deployed first-moment MHE has 16 states and requires a separate timing measurement.

retain `false` as a backward-compatible default, but all 13 release-evaluation batch drivers explicitly set it to `true`. The reported detector results therefore apply to those frozen configurations, not to an unparameterized node launch; without the override, no release decision is evaluated and the controller reaches the conservative UNRESOLVED timeout.

Real-time budget. Both optimizers are acados SQP-RTI (one real-time-iteration SQP step per call, HPIPM QP solver). Table I reports measured wall-clock solve times against the period each loop must meet. The NMPC measurements apply to the deployed 13-state controller: it consumes $\sim 3\%$ of its period at the median and stays under 5% at the 99th percentile. The MHE timing column is retained as a historical measurement of the mass-only 14-state estimator used by the July/August result batches; it *does not* establish the runtime of the final 16-state first-moment estimator. A fresh instrumented timing cohort is required before making an onboard-headroom claim for that final estimator.

VI. RESULTS

A. Event-trigger benefit (estimation and control layers)

Table II (wrench-emulated pure mass step, $n = 6$). The event trigger cuts settling $\sim 4\times$ ($\delta m = -0.3$ kg) / $\sim 2.8\times$ (-0.8 kg); the **no-signal self-trigger recovers the benefit within noise** (~ 0.02 s of the signal version). Control-layer transient peak drops 6% to 14%.

Red line. This experiment reports settling and `pos_err` peak only; overshoot depth is a wrench inertia-mismatch artifact and is not reported here (overshoot conclusions come only from the physical-rigid-body slung experiments).

B. Supporting schedule-tuning ablation

We additionally tested whether the handcrafted transient schedule required numerical optimization. A simulation-

TABLE II

EVENT-TRIGGER VS. FIXED WEIGHTS (WRENCH PURE MASS STEP, $n = 6$).

δm	metric	fixed	signal	no-signal
-0.3	settling [s]	1.20	0.30	0.32
	pos_err peak [m]	0.200	0.187	0.173
-0.8	settling [s]	1.57	0.56	0.57
	pos_err peak [m]	0.545	0.483	0.482




Fig. 3. Mass-estimate timeline (wrench step, $\delta m = -0.5$ kg). The event-triggered estimate reaches the truth band far earlier than fixed weights, at the cost of a transient overshoot below the band before settling—the speed/overshoot trade-off reported honestly.

in-the-loop cross-entropy search embedded M0 in a five-parameter family and evaluated 96 closed-loop episodes. The surviving candidate was then compared with M0 in a pre-registered paired 2×2 factorial ($n = 8$ per cell). Its median settling-time gain was -0.05 s for a 0.2 kg payload and 0.00 s for a 0.3 kg payload (both non-significant); the test had 97 % power against the larger previously observed effect. We therefore retain M0. This ablation supports a deliberately narrow conclusion: for the tested family, *when* to discard stale stages matters more than the precise weight magnitude. It is not a learning contribution of the deployed system.

C. Eventless grasp-release lifecycle

The controller issues the gripper-release command itself, but deliberately does not use that command to clear its payload model: a commanded release is not proof of physical detachment. We compare an event-driven operational reference (arm A), which resets payload geometry and inner-loop scaling with the command, against the proposed interface (arm B), which consumes only healthy, fresh $m/s/J$ frames. Arm B completes a release only after its no-payload confidence remains above 0.90 for 1.0 s. The two arms were paired under three payload and trajectory conditions. A pre-registered pair was voided and re-flown if either arm crashed before release or lacked payload evidence; these exclusions are reported below and are not absorbed into the completion result.

Among valid pairs, arm B completed autonomous release confirmation in 22/22 flights with no observed post-release

TABLE III

EVENTLESS RELEASE COMPLETION FOR ARM B (VALID PAIRED FLIGHTS ONLY). DELAY INCLUDES THE 1.0 s CONFIDENCE HOLD.

condition	n	completed	delay median [IQR]
W3: 0.15 kg, 4 m/s	8	8/8	2.56 s [2.53, 2.85]
W4: 0.30 kg, 4 m/s	6	6/6	3.00 s [2.75, 3.25]
W5: 0.15 kg, 2 m/s	8	8/8	2.70 s [2.46, 2.95]
pooled	22	22/22	2.70 s [2.55, 2.99]

TABLE IV
TARGETED RELEASE-SAFETY ABLATION.

decision rule	flights	premature	confirmed / unresolved
quantity + residual fast path	9	3	—
fresh moment-gated paths	12	0	11 / 1

instability (Table III). The pooled median delay was 2.70 s, including the 1.0 s confidence hold. The two-sided 95 % Clopper-Pearson interval is [84.6 %, 100 %], so this result demonstrates closed-loop feasibility rather than a zero failure rate. In all 22 flights the payload-presence machine made exactly two transitions—one autonomous load declaration and one autonomous empty declaration—with no carrying-phase transition. The published first moment decayed continuously rather than stepping to zero, confirming that recovery used the intended estimator interface rather than the controller’s release command.

a) *Safety boundary*: The 22/22 batch validates the interface once its release evidence is accepted; it does not validate a universally robust detector. We therefore tested three physically distinct evidence sources: payload quantity, directional thrust residual, and collapse of the normalized first moment. A fast branch combining quantity and residual without moment evidence declared release prematurely in 3/9 randomized flights during a 4 m/s figure-eight. Loaded-flight mass estimates could remain below the nominal empty threshold for up to 48 s, while a dynamic thrust excursion reached -1.69 N, exceeding the -1.47 N weight change of the smallest payload. Neither source is therefore reliable on its own in this regime.

We rejected that branch and require every release decision to include fresh, successful moment-collapse evidence. Ambiguous cases enter UNRESOLVED; a timeout does not clear the model or reset adaptive state. In a subsequent 12-flight smoke test, no flight declared release before the command: 11 releases were confirmed with a median detector delay of 1.9 s, and one remained safely unresolved (Table IV). This is the supported safety conclusion; broader detector robustness remains outside the present claim.

The paired performance comparison is conditional on successful attachment and survival to the release command. Of 76 flights, 18 were voided and re-flown: 15 crashed before release and three lacked MHE payload evidence. The available logs cannot distinguish physical capture failure from estimator non-recognition in those three cases. Among valid pairs, steady

TABLE V
THREE-WAY METHOD $\times$ SCENARIO MATRIX ($n = 5$): POS_ERR PEAK [M] / RECOVERY [S].

cell (mass/ecc)	truth	online	l1
0.2 / 0.05	0.806 / 4.28	0.768 / 3.70	0.821 / 4.26
0.2 / 0.10	0.794 / 3.62	0.810 / 3.92	0.812 / 4.24
0.3 / 0.05	0.834 / 3.56	0.837 / 3.72	0.883 / 4.90
0.3 / 0.10	1.011 / 3.64	0.876 / 3.88	0.891 / 4.54

TABLE VI
WIND FALSE-POSITIVE RATE VS. MAGNITUDE ($n = 3$, NO MASS EVENT).

force [N]	0.5	1.0	1.5	2.0	2.5	3.0
vertical FPR	0/3	0/3	3/3	3/3	3/3	3/3
horizontal FPR	—	0/3	—	0/3	—	0/3

tracking RMS differed by at most 0.8 mm between arms; no degradation of engineering significance was observed, although no equivalence test was pre-registered. The measurable cost of refusing the event shortcut is the 2.6 s to 3.0 s median confirmation delay.

D. Baselines and ablations

L1-NMPC three-way matrix (Table V). {truth, online, l1} $\times$ {0.2, 0.3} kg $\times$ {0.05, 0.10} m, $n = 5$, 60/60 runs stable, zero divergence. (i) Transient peaks are statistically indistinguishable across the three methods (0.77 m to 1.01 m, within-cell differences 0.01 m to 0.09 m)—L1 is a competent strong baseline and our method is not worse on the transient. (ii) **L1 recovers 15 % to 35 % slower**; the mechanism is that its compensation low-pass cutoff ω_c is bounded by closed-loop stability (0.5 stable, 1.0 unstable on this platform), a bandwidth wall that explicit estimation does not face. (iii) L1 produces no interpretable m/c_{xy} (cannot update thrust-margin constraints), has no estimation-layer settling notion, and still requires the same class prior for gain scaling. Implementation (footnote): L1 is a translational 3-DoF gradient-law augmentation using T_{phys} as the observed thrust; using the commanded thrust self-excites, and a ground low-thrust gate is required.

Wind false-trigger ablation (Table VI). Self-trigger false-positive rate vs. sustained wind magnitude (link-wrench injection, no mass event, $n = 3$). (i) Sharp step at the 1.5 N threshold: sub-threshold vertical wind never false-triggers, supra-threshold always does—a sustained vertical force is **inherently indistinguishable from a mass change on the thrust-residual channel** (an honest limitation). (ii) **Immune to horizontal wind**: even at $2\times$ threshold (3 N) horizontal wind never false-triggers, because the detector lives on the collective-thrust channel. (iii) Mitigations: the persistence gate rejects gust pulses; real wind is predominantly horizontal/gusty rather than sustained-DC vertical; the threshold is tunable to the noise floor.

VII. LIMITATIONS AND FUTURE WORK

All system-level evidence is presently SITL-only. Hardware migration replaces simulated motor speeds with ESC telemetry

and must re-identify thrust/torque coefficients, noise floors, timing, and gripper compliance before free flight. The vertical arm r_z is unobservable from the thrust-moment channel and remains a weak rack prior; the s_x /pitch channel is noisier than s_y /roll. The deployed diagonal inertia approximation also drops the small $O(r_{xy}^2)$ and yaw-inertia terms.

Release-detector boundary. First-moment augmentation is now part of the default path. The 22/22 paired result shows that the first-moment output can drive complete NMPC recovery without a drop event, but it does not establish distribution-wide detector robustness. The targeted safety ablation identifies the operative boundary: payload quantity and thrust residual can become simultaneously false-empty during dynamic flight, causing premature model clearing in 3/9 flights when combined without moment evidence. The accepted logic therefore requires fresh, healthy moment-collapse evidence and treats timeouts as UNRESOLVED, never as proof of an empty vehicle. After this correction, 12 flights produced no premature declarations, 11 confirmations, and one safe unresolved outcome.

The remaining detector limitation is its moment reference. A running maximum can absorb dynamic estimation excursions above the physically feasible envelope and thereby make an ordinary loaded moment appear collapsed. Future validation will form this reference only from a healthy, low-dynamic, in-envelope interval and freeze it before aggressive flight. Until that revision is validated across broader conditions, our claim is an eventless, fail-safe control interface rather than a fully robust autonomous release detector.

The current software also has a configuration-coupling defect: the historically named `c_xy_est_enable` parameter is the effective master switch for all release evidence. Although every reported release batch explicitly enabled it, the bare default is `false` and therefore falls through to UNRESOLVED. Before deployment, release decision execution must receive its own explicit enable parameter (or be unconditional in the eventless profile), with startup validation that rejects inconsistent configurations.

Selection and pre-release failures. The main performance tables condition on successful attachment and survival to the release command. Eighteen of 76 flights were voided; arm A had 11/38 pre-release crashes and arm B had 4/38, while arm B alone had 3/38 failures of the pre-registered MHE payload-evidence check. The historical logging recorded an attach request but not the simulator joint-state acknowledgement, so those three cases cannot be separated into physical capture failures versus estimator non-recognition. Nine arm A crashes concentrate in W4, leaving only six valid pairs there. Consequently the paper makes no claim that the interface improves grasp, lift, or early figure-eight stability, and the W4 evidence is weaker than W3/W5. Attachment observability and the pre-release attitude instability require separate experiments rather than reinterpretation of the conditional drop result.

A sustained vertical up/downdraft remains intrinsically confounded with a payload change on a collective-thrust-only residual; the existing wind ablation characterizes this limitation but does not remove it.

Simulator fidelity. We record two known deviations. (i) gz-sim Harmonic provides no supported way to change a link’s inertial properties at runtime such that the DART integrator adopts them: `SetComponentData<Inertial>` updates the ECM component only, so introspection tools report the new value while the physics continues to integrate the load-time mass (gz-sim issue #2733). Our wrench mechanism therefore emulates a *translational* mass step by force injection and leaves the rotational inertia unchanged; accordingly it supports only mass-channel claims, and every claim involving inertia, CoM, or eccentricity comes from the physically attached DetachableJoint platform, where the change is genuine (§V). (ii) The eccentric payload is modeled as rigidly attached; a real magnetic gripper admits small compliance and, for a suspended load, pendulum modes this model does not reproduce. Neither deviation affects the sensing-minimal claim, which concerns which signals the estimator consumes rather than how the disturbance is generated.

On learning. Our simulation-in-the-loop optimization of the weight schedule produced no measurable gain (§VI-B), and we identify the binding constraint as the sample economics of high-fidelity evaluation— ~ 85 s per closed-loop episode with no injectable random seed, hence no common random numbers, hence per-candidate noise on the order of the effect being optimized. Approaches that do not share this constraint remain open: supervised targets extracted per-frame from flight data (rather than one scalar loss per episode), or a lower-fidelity differentiable environment for search with SITL reserved for validation. We regard the negative result as delimiting *where* learning can pay off in this pipeline, not as evidence against learned estimator tuning in general.

VIII. CONCLUSION

We presented an eventless payload-adaptive NMPC architecture that uses motor speeds and odometry to update mass, first mass moment, and inertia without consuming attach/drop notifications or treating a release command as proof of physical detachment. Self-triggered MHE window deweighting accelerates abrupt-change recovery, while the atomic, health-gated $m/s/J$ interface and continuous body-rate scheduling connect the estimate safely to the controller. Conditional on successful attachment and pre-release survival, the interface completed release in 22/22 paired SITL flights with no observed post-release instability. It also matched the transient tracking of an L1-NMPC baseline while recovering 15 % to 35 % faster and retaining interpretable physical parameters.

The release-safety ablation defines the claim boundary. Payload quantity and thrust residual became simultaneously false-empty during aggressive flight, so the accepted design requires fresh moment-collapse evidence and preserves ambiguous cases as UNRESOLVED. This policy produced zero premature declarations in the corrected 12-flight smoke, with 11 confirmed releases and one safe unresolved outcome. The present evidence therefore demonstrates that the payload-event signal can be removed from the adaptive control interface; it does not establish a fully robust release detector. A frozen,

robust moment reference, broader SITL validation, and subsequent HITL and flight testing are the next steps.

REFERENCES

- [1] D. Hanover, P. Foehn, S. Sun, E. Kaufmann, and D. Scaramuzza, “Performance, precision, and payloads: Adaptive nonlinear MPC for quadrotors,” *IEEE Robotics and Automation Letters*, vol. 7, no. 2, pp. 690–697, 2022.
- [2] E. J. J. Smeur, Q. Chu, and G. C. H. E. de Croon, “Adaptive incremental nonlinear dynamic inversion for attitude control of micro air vehicles,” *Journal of Guidance, Control, and Dynamics*, vol. 39, no. 3, pp. 450–461, 2016.
- [3] J. Svacha, J. Paulos, G. Loianno, and V. Kumar, “IMU-based inertia estimation for a quadrotor using Newton–Euler dynamics,” *IEEE Robotics and Automation Letters*, vol. 5, no. 3, pp. 3861–3867, 2020.
- [4] V. Wüest, V. Kumar, and G. Loianno, “Online estimation of geometric and inertia parameters for multirotor aerial vehicles,” in *IEEE International Conference on Robotics and Automation (ICRA)*, 2019, pp. 1884–1890.
- [5] C. Böhm, M. Scheiber, and S. Weiss, “Filter-based online system-parameter estimation for multicopter UAVs,” in *Robotics: Science and Systems (RSS)*, 2021, paper 087, DOI 10.15607/RSS.2021.XVII.087.
- [6] B. Wang, Z. Ma, S. Lai, and L. Zhao, “Neural moving horizon estimation for robust flight control,” *IEEE Transactions on Robotics*, vol. 40, pp. 639–659, 2024, open source: RCL-NUS/NeuroMHE.